

A gentle introduction to Bayesian inference, MCMC, and Stan

Matt Espe

February 10, 2017

Data Science Initiative, University of California - Davis

Outline

Bayesian inference (a crash course)

Markov Chain Monte Carlo

Comparing MCMC algorithms

Troublesome posteriors

Some caution required

Conclusion

Demonstration

Bayesian inference (a crash course)

Bayes Theorem

$$P(\theta | X) = \frac{P(X | \theta) P(\theta)}{P(X)}$$

$P(\theta | X)$ is the posterior,

$P(X | \theta)$ is the likelihood,

$P(\theta)$ is the prior,

$P(X)$ is the support.

$$P(\theta \mid X) \propto P(X \mid \theta) P(\theta)$$

Why Bayes?

It presents a coherent method to combine information from multiple sources.

It is flexible.

However, it will not save you from yourself!

The issue with Bayes

For simple problems, we can analytically derive the posterior.

More complicated models might not be analytically tractable.

But, we can sample from $P(\theta | X)$ even without knowing the analytical solution!

Using samples from $P(\theta | X)$, we can calculate MCMC estimators.

Markov Chain Monte Carlo

Markov Chain Monte Carlo (MCMC) is a method to sample from $P(\theta \mid X)$.

Many different algorithms with different strengths and weaknesses.

MCMC estimators will converge to true expectations (eventually).

MCMC algorithms are not magic.

- The good: You can write custom samplers, develop highly flexible models, and do generally cool stuff.
- The bad: They cannot overcome issues with model specification, data collection or experimental design, etc.
- The ugly: You never know if they are working properly, only if they are not obviously broken.

Random Walk Metropolis Hastings

```
RWmetroNorm <- function(x, prior_mu, prior_sd, known_sd = 1,
                        step_size = 0.1, iter = 100)
{
  samples <- numeric(iter)
  samples[1] <- rnorm(1, prior_mu, prior_sd)


  for(i in 2:iter){
    prop <- rnorm(1, samples[i - 1], step_size)

    cur_ll <- sum(dnorm(x, samples[i - 1], sd = known_sd, log = TRUE)) +
      dnorm(samples[i - 1], prior_mu, prior_sd, log = TRUE)
    prop_ll <- sum(dnorm(x, prop, sd = known_sd, log = TRUE)) +
      dnorm(prop, prior_mu, prior_sd, log = TRUE)

    jump_prob <- min(exp(prop_ll - cur_ll), 1)


    if(runif(1) < jump_prob){
      samples[i] <- prop
    } else {
      samples[i] <- samples[i - 1]
    }
  }
  return(samples)
}
```

Random Walk MH (cont.)



N = 100000 Bandwidth = 0.08503

The 'fuzzy caterpillar'



JAGS: Just Another Gibbs Sampler

A popular MCMC sampler with cross-platform support.

```
mod.text <- '  
model{  
  for( i in 1:N ){  
    x[i] ~ dnorm(mu, (1/known_sd ^2))  
  }  
  mu ~ dnorm(prior_mu, (1/prior_sd^2))  
}'
```

Comparing MCMC algorithms

Comparing MCMC algorithms

Which one is faster?

RW_time

| | | | |
|----|-------|--------|---------|
| ## | user | system | elapsed |
| ## | 1.453 | 0.067 | 1.518 |

JAGS_time

| | | | |
|----|-------|--------|---------|
| ## | user | system | elapsed |
| ## | 0.053 | 0.000 | 0.051 |

Is faster what we want?

Not necessarily - how do we measure speed?

Comparing MCMC (cont.)

Effective size: The equivalent number of independent draws.

```
effectiveSize(ans)
```

```
##      var1
```

```
## 4782.68
```

```
effectiveSize(jags.ans)
```

```
##      mu
```

```
## 1e+05
```

Comparing MCMC (cont.)

Effective size/time is often what we care about (efficiency).

As the effective number of samples increases, the MCMC estimators will converge to the true estimator.

Unless high precision (low MCMC error) is needed, 1000 effective samples is sufficient for most applications.

Troublesome posteriors


Let's break things!

```
x2 <- rnorm(25, x, sd = 0.1)
y <- rnorm(25, x * 2, 0.5)


mod2.text <- '
model{
  for( i in 1:N ){
    y[i] ~ dnorm(beta1 * x[i] + beta2 * x2[i],
                 (1/known_sd ^2))
  }
  beta1 ~ dnorm(0, 1)
  beta2 ~ dnorm(0, 1)
}'
```

Difficult posteriors


Trace of beta1




Trace of beta2



(Enter Stan)



Why does this break?



What is the impact of this?

```
effectiveSize(jags.ans2)
```

```
##      beta1      beta2
```

```
## 14.52412 14.50943
```

```
summary(stan.mod)[1:2,9]
```

```
##      x      x2
```

```
## 940 940
```

MCMC efficiency revisited

JAGS is more efficient due to its speed, but requires many more draws.

```
effectiveSize(jags.ans2)/jags_time['elapsed']
```

```
##      beta1      beta2  
## 453.8787 453.4197
```

```
summary(stan.mod)[1:2,9]/stan_time['elapsed']
```

```
##          x          x2  
## 57.7431 57.7431
```

Typical JAGS solution

Sample a ton, then throw out all but every n draw.

100,000 iterations with a 50 draw thin:

```
effectiveSize(jags.ans2)
```

```
##      beta1      beta2
```

```
## 1800.337 1802.250
```

A more difficult model to fit

$$y = \alpha + \beta_{site} + \beta_{cultivar} + \beta_{year} + \beta_{year:site}$$

$N = 4087$

$n \text{ cult} = 14$

$n \text{ year} = 21$

$n \text{ site} = 25$

$n \text{ site:year} = 168$

```
m3 <- 'model{
  for( i in 1:N ){
    mu[i] = alpha + b_site[site[i]] +
      b_cult[cult[i]] + b_year[year[i]] +
      b_siteyr[siteyr[i]]
    yield[i] ~ dnorm(mu[i],
      (1/sigma ^2))
  }
  alpha ~ dnorm(0, 1)
  for(i in 1:n_site){
    b_site[i] ~ dnorm(0, (1/pow(tau[1],2)))}
  for(i in 1:n_cult){
    b_cult[i] ~ dnorm(0, (1/pow(tau[2],2)))}
  for(i in 1:n_year){
    b_year[i] ~ dnorm(0, (1/pow(tau[3],2)))}
  for(i in 1:n_yr_site){
    b_siteyr[i] ~ dnorm(0, (1/pow(tau[4],2)))}
  for(i in 1:4){
    tau[i] ~ dunif(0, 10)}
  sigma ~ dunif(0, 10)
}'
```

```
effectiveSize(jags.fit)[1:5]/jags.time2["elapsed"]
```

```
## b_cult[1] b_cult[2] b_cult[3] b_cult[4] b_cult[5]  
## 2.179756 2.318169 3.774028 4.925949 1.893628
```

```
summary(stan.fit, regex_pars = " id")[1:5, 9]/stan.time2["elapsed"]
```

```
## b[(Intercept) id:M103] b[(Intercept) id:M104] b[(Intercept) id:M105]  
## 3.421933 3.421933 3.421933  
## b[(Intercept) id:M201] b[(Intercept) id:M202]  
## 3.421933 3.421933
```

Some caution required

\hat{R} compares within and between chain variability

N_{eff} measures the effective draws a sample contains

N_{eff}/s measures efficiency

$N_{\text{eff}}/\text{draw}$ is a different measure of efficiency

Traceplots show how well chains mix (i.e., traverse the posterior)

WAIC, PS-loo, Fraction of Bayesian missing information, etc...

Warning!

MCMC is a powerful tool but:

OK diagnostics are a sign that nothing has gone obviously wrong, not a sign that things have gone right.

On that note...

Always run multiple chains, starting from different initial values. A good rule of thumb is minimum of 4 chains, more if possible.

(I typically run 8-16 chains)

More chains increases the chances that one will find a pathological region in the posterior, numeric instability, etc.

Do not just remove a poorly sampling chain! It is a sign something is wrong!

Conclusion

So why Stan?

Stan was built specifically to handle:

1. Difficult posteriors
2. High dimensional problems

Additional benefits

Stan is flexible*

Under active development

Large user base

Cool add-on tools

Conclusions

Bayesian inference for most non-trivial problems requires some way to sample from the posterior distribution, $P(\theta \mid X)$.

Samples can be generated using MCMC algorithms

Algorithms differ in their speed and efficiency

Although flexible and powerful, lots can go wrong

Demonstration
